

Outline

It is always good to create helper functions to assist you in making a model basically breaking it down into puzzle pieces.

- 1 Initialize the parameters for a two-layer network and for an L-layer neural network
2. Implement the forward propagation module
 - Complete the LINEAR part of a layer's forward propagation step (resulting in $Z[l]$).
 - The ACTIVATION function is provided for you (relu/sigmoid)
 - Combine the previous two steps into a new [LINEAR->ACTIVATION] forward function.
 - Stack the [LINEAR->RELU] forward function L-1 time (for layers 1 through L-1) and add a [LINEAR->SIGMOID] at the end (for the final layer L). This gives you a new `L_model_forward` function.
3. Compute the loss
4. Implement the backward propagation module (denoted in red in the figure below)
 - Complete the LINEAR part of a layer's backward propagation step
 - The gradient of the ACTIVATION function is provided for you (relu_backward/sigmoid_backward)
 - Combine the previous two steps into a new [LINEAR->ACTIVATION] backward function
 - Stack [LINEAR->RELU] backward L-1 times and add [LINEAR->SIGMOID] backward in a new `L_model_backward` function
5. Finally, update the parameters

For every forward function, there is a corresponding backward function. This is why at every step of your forward module you will be storing some values in a cache. These cached values are useful for computing gradients.

In the backpropagation module, you can then use the cache to calculate the gradients.

1. Initialize the parameters for a two-layer network and for an L-layer neural network

2 Layer Neural Network function:

def initialize_parameters(n_x, n_h, n_y):

"""

Argument:

n_x -- size of the input layer

n_h -- size of the hidden layer

n_y -- size of the output layer

Returns:

parameters -- python dictionary containing your parameters:

W1 -- weight matrix of shape (n_h, n_x)

b1 -- bias vector of shape (n_h, 1)

W2 -- weight matrix of shape (n_y, n_h)

b2 -- bias vector of shape (n_y, 1)

"""

np.random.seed(1)

W1 = np.random.randn(n_h, n_x) * .01 # Small random values

W2 = np.random.randn(n_y, n_h) * .01

b1 = np.zeros((n_h, 1)) # Np array of zeros

b2 = np.zeros((n_y, 1))

parameters = {"W1": W1,

"b1": b1,

"W2": W2,

"b2": b2}

return parameters

L Layer Neural Network:

```
def initialize_parameters_deep(layer_dims):
```

```
    """
```

Arguments:

layer_dims -- python array (list) containing the dimensions of each layer in our network

Returns:

parameters -- python dictionary containing your parameters "W1", "b1", ..., "WL", "bL":

 Wl -- weight matrix of shape (layer_dims[l], layer_dims[l-1])

 bl -- bias vector of shape (layer_dims[l], 1)

```
    """
```

```
np.random.seed(3)
```

```
parameters = {}
```

```
L = len(layer_dims) # number of layers in the network
```

```
for l in range(1, L):
```

```
    # So per Layer
```

```
    parameters["W" + str(l)] = np.random.randn(layer_dims[l], layer_dims[l-1]) * 0.01
```

```
    parameters["b" + str(l)] = np.zeros((layer_dims[l], 1))
```

```
    assert(parameters["W" + str(l)].shape == (layer_dims[l], layer_dims[l - 1]))
```

```
    assert(parameters["b" + str(l)].shape == (layer_dims[l], 1))
```

```
return parameters
```

2. Forward Propagation

- **Complete the LINEAR part of a layer's forward propagation step (resulting in $Z[l]$).**

```
def linear_forward(A, W, b):
```

```
    """
```

Implement the linear part of a layer's forward propagation.

Arguments:

A -- activations from previous layer (or input data): (size of previous layer, number of examples)

W -- weights matrix: numpy array of shape (size of current layer, size of previous layer)

b -- bias vector, numpy array of shape (size of the current layer, 1)

Returns:

Z -- the input of the activation function, also called pre-activation parameter

cache -- a python tuple containing "A", "W" and "b" ; stored for computing the backward pass efficiently

```
    """
```

```
    Z = np.dot(W,A) + b
```

```
    cache = (A, W, b) #Storing for backpropagation
```

```
    return Z, cache
```

- **The ACTIVATION function is provided for you (relu/sigmoid)**
- **Combine the previous two steps into a new [LINEAR->ACTIVATION] forward function.**

```
def linear_activation_forward(A_prev, W, b, activation):
```

```
"""
```

Implement the forward propagation for the LINEAR->ACTIVATION layer

Arguments:

A_prev -- activations from previous layer (or input data): (size of previous layer, number of examples)

W -- weights matrix: numpy array of shape (size of current layer, size of previous layer)

b -- bias vector, numpy array of shape (size of the current layer, 1)

activation -- the activation to be used in this layer, stored as a text string: "sigmoid" or "relu"

Returns:

A -- the output of the activation function, also called the post-activation value

cache -- a python tuple containing "linear_cache" and "activation_cache";

stored for computing the backward pass efficiently

```
"""
```

```
Z, linear_cache = linear_forward(A_prev,W,b) # Linear part
```

```
if activation == "sigmoid":
```

```
    A, activation_cache = sigmoid(Z)
```

```
elif activation == "relu":
```

```
    A, activation_cache = relu(Z)
```

```
cache = (linear_cache, activation_cache)
```

```
return A, cache
```

- **Stack the [LINEAR->RELU] forward function L-1 time (for layers 1 through L-1) and add a [LINEAR->SIGMOID] at the end (for the final layer L). This gives you a new L_model_forward function.**

```
def L_model_forward(X, parameters):
```

```
    """
```

Implement forward propagation for the [LINEAR->RELU]*(L-1)->LINEAR->SIGMOID computation

Arguments:

X -- data, numpy array of shape (input size, number of examples)

parameters -- output of initialize_parameters_deep()

Returns:

AL -- activation value from the output (last) layer

caches -- list of caches containing:

every cache of linear_activation_forward() (there are L of them, indexed from 0 to L-1)

```
    """
```

```
    caches = []
```

```
    A = X
```

```
    L = len(parameters) // 2 # number of layers in the neural network
```

```
    # Implement [LINEAR -> RELU]*(L-1). Add "cache" to the "caches" list.
```

```
    # The for loop starts at 1 because layer 0 is the input
```

```
    for l in range(1, L):
```

```
        A_prev = A
```

```
        A, cache = linear_activation_forward(A_prev, parameters["W" + str(l)], parameters["b" + str(l)],
activation="relu")
```

```
        caches.append(cache)
```

```

# Implement LINEAR -> SIGMOID. Add "cache" to the "caches" list.

AL, cache = linear_activation_forward(A, parameters["W" + str(L)], parameters["b" + str(L)],
activation="sigmoid")

caches.append(cache)


return AL, caches

```

3. Compute Cost

```
def compute_cost(AL, Y):
```

```
    """
```

Implement the cost function defined by equation.

Arguments:

AL -- probability vector corresponding to your label predictions, shape (1, number of examples)

Y -- true "label" vector (for example: containing 0 if non-cat, 1 if cat), shape (1, number of examples)

Returns:

cost -- cross-entropy cost

```
    """
```

```
    m = Y.shape[1]
```

```
    # Compute loss from aL and y.
```

```
    cost = -(1/m) * np.sum(Y * np.log(AL) + (1-Y) * np.log(1-AL))
```

```
    cost = np.squeeze(cost) # To make sure your cost's shape is what we expect (e.g. this turns
[[17]] into 17).
```

```
    return cost
```

4. Backward Propagation

■ Complete the LINEAR part of a layer's backward propagation step

```
def linear_backward(dZ, cache):
```

```
    """
```

Implement the linear portion of backward propagation for a single layer (layer l)

Arguments:

dZ -- Gradient of the cost with respect to the linear output (of current layer l)

cache -- tuple of values (A_prev, W, b) coming from the forward propagation in the current layer

Returns:

dA_prev -- Gradient of the cost with respect to the activation (of the previous layer l-1), same shape as A_prev

dW -- Gradient of the cost with respect to W (current layer l), same shape as W

db -- Gradient of the cost with respect to b (current layer l), same shape as b

```
    """
```

```
    A_prev, W, b = cache
```

```
    m = A_prev.shape[1]
```

```
    dW = (1/m) * np.dot(dZ, A_prev.T)
```

```
    db = (1/m) * np.sum(dZ, axis=1, keepdims=True)
```

```
    dA_prev = np.dot(W.T, dZ)
```

```
    return dA_prev, dW, db
```


- The gradient of the **ACTIVATION** function is provided for you(`relu_backward/sigmoid_backward`)
- Combine the previous two steps into a new **[LINEAR->ACTIVATION] backward function**

```
def linear_activation_backward(dA, cache, activation):
```

```
    """
```

Implement the backward propagation for the LINEAR->ACTIVATION layer.

Arguments:

`dA` -- post-activation gradient for current layer `l`

`cache` -- tuple of values (`linear_cache`, `activation_cache`) we store for computing backward propagation efficiently

`activation` -- the activation to be used in this layer, stored as a text string: "sigmoid" or "relu"

Returns:

`dA_prev` -- Gradient of the cost with respect to the activation (of the previous layer `l-1`), same shape as `A_prev`

`dW` -- Gradient of the cost with respect to `W` (current layer `l`), same shape as `W`

`db` -- Gradient of the cost with respect to `b` (current layer `l`), same shape as `b`

```
    """
```

```
    linear_cache, activation_cache = cache
```

```
    if activation == "relu":
```

```
        dZ = relu_backward(dA, activation_cache)
```

```
    elif activation == "sigmoid":
```

```
        dZ = sigmoid_backward(dA, activation_cache)
```

```
    dA_prev, dW, db = linear_backward(dZ, linear_cache)
```

```
    return dA_prev, dW, db
```

- **Stack [LINEAR->RELU] backward L-1 times and add [LINEAR->SIGMOID] backward in a new L_model_backward function**

```
def L_model_backward(AL, Y, caches):
```

```
    """
```

Implement the backward propagation for the [LINEAR->RELU] * (L-1) -> LINEAR -> SIGMOID group

Arguments:

AL -- probability vector, output of the forward propagation (L_model_forward())

Y -- true "label" vector (containing 0 if non-cat, 1 if cat)

caches -- list of caches containing:

every cache of linear_activation_forward() with "relu" (it's caches[l], for l in range(L-1) i.e l = 0...L-2)

the cache of linear_activation_forward() with "sigmoid" (it's caches[L-1])

Returns:

grads -- A dictionary with the gradients

```
    grads["dA" + str(l)] = ...
```

```
    grads["dW" + str(l)] = ...
```

```
    grads["db" + str(l)] = ...
```

```
    """
```

```
    grads = {}
```

```
    L = len(caches) # the number of layers
```

```
    m = AL.shape[1]
```

```
    Y = Y.reshape(AL.shape) # after this line, Y is the same shape as AL
```

```
    # Initializing the backpropagation
```

```
    # Initial Gradient of Cost with respect to AL
```

```
dAL = - (np.divide(Y, AL) - np.divide(1 - Y, 1 - AL))
```

```
# Compute gradients for the last layer (Sigmoid Activation)
```

```
current_cache = caches[L-1]
```

```
dA_prev, dW, db = linear_activation_backward(dAL, current_cache, activation="sigmoid")
```

```
grads["dA" + str(L-1)] = dA_prev
```

```
grads["dW" + str(L)] = dW
```

```
grads["db" + str(L)] = db
```

```
# Loop through the remaining layers in reverse order (ReLU Activation)
```

```
# Loop from l=L-2 to l=0
```

```
for l in reversed(range(L-1)):
```

```
    # lth layer: (RELU -> LINEAR) gradients.
```

```
    # Inputs: "grads["dA" + str(l + 1)], current_cache". Outputs: "grads["dA" + str(l)] ,  
grads["dW" + str(l + 1)] , grads["db" + str(l + 1)]
```

```
    current_cache = caches[l]
```

```
    dA_prev, dW, db = linear_activation_backward(grads["dA" + str(l + 1)], current_cache,  
activation="relu")
```

```
grads["dA" + str(l)] = dA_prev
```

```
grads["dW" + str(l + 1)] = dW
```

```
grads["db" + str(l + 1)] = db
```

```
return grads
```

5. Update Parameters

```
def update_parameters(params, grads, learning_rate):
```

```
    """
```

```
    Update parameters using gradient descent
```

Arguments:

params -- python dictionary containing your parameters

grads -- python dictionary containing your gradients, output of L_model_backward

Returns:

parameters -- python dictionary containing your updated parameters

```
        parameters["W" + str(l)] = ...
```

```
        parameters["b" + str(l)] = ...
```

```
    """
```

```
    parameters = copy.deepcopy(params)
```

```
    L = len(parameters) // 2 # number of layers in the neural network
```

```
    # Update rule for each parameter. Use a for loop.
```

```
    for l in range(L):
```

```
        parameters["W" + str(l+1)] -= learning_rate * grads["dW" + str(l+1)]
```

```
        parameters["b" + str(l+1)] -= learning_rate * grads["db" + str(l+1)]
```

```
    return parameters
```

THE MODELS

1.

```
def two_layer_model(X, Y, layers_dims, learning_rate = 0.0075, num_iterations = 3000,
print_cost=False):
```

```
    """
```

Implements a two-layer neural network: LINEAR->RELU->LINEAR->SIGMOID.

Arguments:

X -- input data, of shape (n_x, number of examples)

Y -- true "label" vector (containing 1 if cat, 0 if non-cat), of shape (1, number of examples)

layers_dims -- dimensions of the layers (n_x, n_h, n_y)

num_iterations -- number of iterations of the optimization loop

learning_rate -- learning rate of the gradient descent update rule

print_cost -- If set to True, this will print the cost every 100 iterations

Returns:

parameters -- a dictionary containing W1, W2, b1, and b2

```
    """
```

```
    np.random.seed(1)
```

```
    grads = {}
```

```
    costs = []           # to keep track of the cost
```

```
    m = X.shape[1]       # number of examples
```

```
    (n_x, n_h, n_y) = layers_dims
```

```

# Initialize parameters dictionary, by calling one of the functions you'd previously implemented
parameters = initialize_parameters(n_x, n_h, n_y)

# Get W1, b1, W2 and b2 from the dictionary parameters.
W1 = parameters["W1"]
b1 = parameters["b1"]
W2 = parameters["W2"]
b2 = parameters["b2"]

# Loop (gradient descent)

for i in range(0, num_iterations):

    # Forward propagation: LINEAR -> RELU -> LINEAR -> SIGMOID. Inputs: "X, W1, b1, W2,
    b2". Output: "A1, cache1, A2, cache2".

    A1, cache1 = linear_activation_forward(X, W1, b1, activation="relu")
    A2, cache2 = linear_activation_forward(A1, W2, b2, activation="sigmoid")

    # Compute cost

    cost = compute_cost(A2, Y)

    # Initializing backward propagation

    dA2 = - (np.divide(Y, A2) - np.divide(1 - Y, 1 - A2))

    # Backward propagation. Inputs: "dA2, cache2, cache1". Outputs: "dA1, dW2, db2; also
    dA0 (not used), dW1, db1".

```

```

dA1, dW2, db2 = linear_activation_backward(dA2, cache2, activation="sigmoid")
dA0, dW1, db1 = linear_activation_backward(dA1, cache1, activation="relu")

# Set grads['dW1'] to dW1, grads['db1'] to db1, grads['dW2'] to dW2, grads['db2'] to db2
grads['dW1'] = dW1
grads['db1'] = db1
grads['dW2'] = dW2
grads['db2'] = db2

# Update parameters.
parameters = update_parameters(parameters, grads, learning_rate)

# Retrieve W1, b1, W2, b2 from parameters
W1 = parameters["W1"]
b1 = parameters["b1"]
W2 = parameters["W2"]
b2 = parameters["b2"]

# Print the cost every 100 iterations
if print_cost and i % 100 == 0 or i == num_iterations - 1:
    print("Cost after iteration {}: {}".format(i, np.squeeze(cost)))
if i % 100 == 0 or i == num_iterations:
    costs.append(cost)

return parameters, costs

```

```
def plot_costs(costs, learning_rate=0.0075):
    plt.plot(np.squeeze(costs))
    plt.ylabel('cost')
    plt.xlabel('iterations (per hundreds)')
    plt.title("Learning rate =" + str(learning_rate))
    plt.show()
```

2.

```
def L_layer_model(X, Y, layers_dims, learning_rate = 0.0075, num_iterations = 3000,
print_cost=False):
```

```
    """
```

Implements a L-layer neural network: [LINEAR->RELU]*(L-1)->LINEAR->SIGMOID.

Arguments:

X -- input data, of shape (n_x, number of examples)

Y -- true "label" vector (containing 1 if cat, 0 if non-cat), of shape (1, number of examples)

layers_dims -- list containing the input size and each layer size, of length (number of layers + 1).

learning_rate -- learning rate of the gradient descent update rule

num_iterations -- number of iterations of the optimization loop

print_cost -- if True, it prints the cost every 100 steps

Returns:

parameters -- parameters learnt by the model. They can then be used to predict.


```
"""
```

```
np.random.seed(1)
```

```
costs = []          # keep track of cost
```

```
# Parameters initialization.
```

```
parameters = initialize_parameters_deep(layers_dims)
```

```
# Loop (gradient descent)
```

```
for i in range(0, num_iterations):
```

```
    # Forward propagation: [LINEAR -> RELU]*(L-1) -> LINEAR -> SIGMOID.
```

```
    AL, caches = L_model_forward(X, parameters)
```

```
    # Compute cost.
```

```
    cost = compute_cost(AL, Y)
```

```
    # Backward propagation.
```

```
    grads = L_model_backward(AL, Y, caches)
```

```
    # Update parameters.
```

```
    parameters = update_parameters(parameters, grads, learning_rate)
```

```
# YOUR CODE ENDS HERE
```

```
# Print the cost every 100 iterations

if print_cost and i % 100 == 0 or i == num_iterations - 1:
    print("Cost after iteration {}: {}".format(i, np.squeeze(cost)))
if i % 100 == 0 or i == num_iterations:
    costs.append(cost)

return parameters, costs
```